

# 1 Predicting High School Student Performance

Mary Nathalie Dela Cruz & Joshua Dela Cruz

The data set contains 30 descriptors on 600+ students from two Portuguese schools. The goal is to predict G1, G2, G3, which are the 1st, 2nd, and 3rd period grades, respectively—a regression problem. The accompanying paper to this data set can be found in: <http://www3.dsi.uminho.pt/pcortez/student.pdf> (<http://www3.dsi.uminho.pt/pcortez/student.pdf>)

## Methodology:

1. Load the Student Performance data set from the UCI Repository. For this problem, we will only deal with the Math scores.
2. Make the necessary encoding for categorical inputs.
3. Split the data into 80% Training and 20% Testing
4. Run any AutoML procedure (either LazyPredict, Optuna, TPOT, or Auto-sklearn) to predict the G3 score using the 30 descriptors and the G1 and G2 scores as input features (32 features all in all). You may limit your search to only a few ML models, especially if you use Optuna.
5. Report the R2 metric on all models that were tried.
6. Based on your results, perform an explainability analysis on the best model using Shapley values.
7. Report the summary plot of the most influential descriptors, then write a discussion on your analysis.

```
In [33]: from ucimlrepo import fetch_ucirepo

import matplotlib.pyplot as plt
import seaborn as sns

import pandas as pd
import numpy as np

import zipfile
import urllib.request

from sklearn.preprocessing import StandardScaler, OrdinalEncoder, OneHotEncoder, LabelEncoder, MinMaxScaler
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split

from lazypredict.Supervised import LazyRegressor

import shap

pd.set_option('display.max_colwidth', None)
```

## ▼ 1.0.1 Exploratory Data Analysis

```
In [34]: student_performance = fetch_ucirepo(id=320)

# Show metadata
metadata = student_performance.metadata
for key, value in metadata.items():
    print(key, ":", value)

uci_id : 320
name : Student Performance
repository_url : https://archive.ics.uci.edu/dataset/320/student+performance (https://archive.ics.uci.edu/dataset/320/student+performance)
data_url : https://archive.ics.uci.edu/static/public/320/data.csv (https://archive.ics.uci.edu/static/public/320/data.csv)
abstract : Predict student performance in secondary education (high school).
area : Social Science
tasks : ['Classification', 'Regression']
characteristics : ['Multivariate']
num_instances : 649
num_features : 30
feature_types : ['Integer']
demographics : ['Sex', 'Age', 'Other', 'Education Level', 'Occupation']
target_col : ['G1', 'G2', 'G3']
index_col : None
has_missing_values : no
missing_values_symbol : None
year_of_dataset_creation : 2008
last_updated : Feb 24, 2024
```

```
In [35]: # Show additional information about dataset
additional_info = metadata["additional_info"]
for key, value in additional_info.items():
    print(key, ":", value)
```

summary : This data approach student achievement in secondary education of two Portuguese schools. The data attributes include student grades, demographic, social and school related features) and it was collected by using school reports and questionnaires. Two datasets are provided regarding the performance in two distinct subjects: Mathematics (mat) and Portuguese language (por). In [Cortez and Silva, 2008], the two datasets were modeled under binary/five-level classification and regression tasks. Important note: the target attribute G3 has a strong correlation with attributes G2 and G1. This occurs because G3 is the final year grade (issued at the 3rd period), while G1 and G2 correspond to the 1st and 2nd period grades. It is more difficult to predict G3 without G2 and G1, but such prediction is much more useful (see paper source for more details).

purpose : None  
funded\_by : None  
instances\_represent : None  
recommended\_data\_splits : None  
sensitive\_data : None  
preprocessing\_description : None  
variable\_info : # Attributes for both student-mat.csv (Math course) and student-por.csv (Portuguese language course) datasets:  
1 school - student's school (binary: 'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira)  
2 sex - student's sex (binary: 'F' - female or 'M' - male)  
3 age - student's age (numeric: from 15 to 22)  
4 address - student's home address type (binary: 'U' - urban or 'R' - rural)  
5 famsize - family size (binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3)  
6 Pstatus - parent's cohabitation status (binary: 'T' - living together or 'A' - apart)  
7 Medu - mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)  
8 Fedu - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)  
9 Mjob - mother's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at\_home' or 'other')  
10 Fjob - father's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at\_home' or 'other')  
11 reason - reason to choose this school (nominal: close to 'home', school 'reputation', 'course' preference or 'other')  
12 guardian - student's guardian (nominal: 'mother', 'father' or 'other')  
13 traveltime - home to school travel time (numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour)  
14 studytime - weekly study time (numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours)  
15 failures - number of past class failures (numeric: n if 1<=n<3, else 4)  
16 schoolsup - extra educational support (binary: yes or no)  
17 famsup - family educational support (binary: yes or no)  
18 paid - extra paid classes within the course subject (Math or Portuguese) (binary: yes or no)  
19 activities - extra-curricular activities (binary: yes or no)  
20 nursery - attended nursery school (binary: yes or no)  
21 higher - wants to take higher education (binary: yes or no)  
22 internet - Internet access at home (binary: yes or no)  
23 romantic - with a romantic relationship (binary: yes or no)  
24 famrel - quality of family relationships (numeric: from 1 - very bad to 5 - excellent)  
25 freetime - free time after school (numeric: from 1 - very low to 5 - very high)  
26 goout - going out with friends (numeric: from 1 - very low to 5 - very high)  
27 Dalc - workday alcohol consumption (numeric: from 1 - very low to 5 - very high)  
28 Walc - weekend alcohol consumption (numeric: from 1 - very low to 5 - very high)  
29 health - current health status (numeric: from 1 - very bad to 5 - very good)  
30 absences - number of school absences (numeric: from 0 to 93)  
  
# these grades are related with the course subject, Math or Portuguese:  
31 G1 - first period grade (numeric: from 0 to 20)  
31 G2 - second period grade (numeric: from 0 to 20)  
32 G3 - final grade (numeric: from 0 to 20, output target)  
citation : None

```
In [36]: # Show scientific article information
intro_paper = metadata["intro_paper"]
for key, value in intro_paper.items():
    print(key, ":", value)
```

title : Using data mining to predict secondary school student performance  
authors : P. Cortez, A. M. G. Silva  
published\_in : Proceedings of 5th Annual Future Business Technology Conference  
year : 2008  
url : <https://www.semanticscholar.org/paper/61d468d5254730bbecf822c6b60d7d6595d9889c> (<https://www.semanticscholar.org/paper/61d468d5254730bbecf822c6b60d7d6595d9889c>)  
doi : None

```
In [37]: # Show variable information
variables = student_performance.variables
display(variables)
```

	name	role	type	demographic	description	units	missing_values
0	school	Feature	Categorical	None	student's school (binary: 'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira)	None	no
1	sex	Feature	Binary	Sex	student's sex (binary: 'F' - female or 'M' - male)	None	no
2	age	Feature	Integer	Age	student's age (numeric: from 15 to 22)	None	no
3	address	Feature	Categorical	None	student's home address type (binary: 'U' - urban or 'R' - rural)	None	no
4	famsize	Feature	Categorical	Other	family size (binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3)	None	no
5	Pstatus	Feature	Categorical	Other	parent's cohabitation status (binary: 'T' - living together or 'A' - apart)	None	no
6	Medu	Feature	Integer	Education Level	mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)	None	no
7	Fedu	Feature	Integer	Education Level	father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)	None	no

```
In [38]: # We just want the Mathematics dataset so we download from the URL
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/00320/student.zip"
filehandle, _ = urllib.request.urlretrieve(url)
zip_file_object = zipfile.ZipFile(filehandle, 'r')
zip_file_object.extractall()

df = pd.read_csv('student-mat.csv', sep=";")
columns = df.columns
df.head()
```

```
Out[38]:
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	famrel	freetime	goout	Dalc	Walc	health	a
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	4	3	4	1	1	3	
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	5	3	3	1	1	3	
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	4	3	2	2	3	3	
3	GP	F	15	U	GT3	T	4	2	health	services	...	3	2	2	1	1	5	
4	GP	F	16	U	GT3	T	3	3	other	other	...	4	3	2	1	2	5	

5 rows × 33 columns

In [39]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
#   Column          Non-Null Count  Dtype
---  -
0   school          395 non-null   object
1   sex             395 non-null   object
2   age            395 non-null   int64
3   address         395 non-null   object
4   famsize         395 non-null   object
5   Pstatus        395 non-null   object
6   Medu           395 non-null   int64
7   Fedu           395 non-null   int64
8   Mjob           395 non-null   object
9   Fjob           395 non-null   object
10  reason         395 non-null   object
11  guardian       395 non-null   object
12  traveltime     395 non-null   int64
13  studytime      395 non-null   int64
14  failures       395 non-null   int64
15  schoolsup      395 non-null   object
16  famsup         395 non-null   object
17  paid           395 non-null   object
18  activities     395 non-null   object
19  nursery       395 non-null   object
20  higher         395 non-null   object
21  internet       395 non-null   object
22  romantic       395 non-null   object
23  famrel        395 non-null   int64
24  freetime      395 non-null   int64
25  goout         395 non-null   int64
26  Dalc          395 non-null   int64
27  Walc          395 non-null   int64
28  health        395 non-null   int64
29  absences      395 non-null   int64
30  G1            395 non-null   int64
31  G2            395 non-null   int64
32  G3            395 non-null   int64
dtypes: int64(16), object(17)
memory usage: 102.0+ KB
```

```
In [40]: # Check for missing values
print(f"NO. OF MISSING VALLUES IN DF: {df.isnull().sum().sum()}")
# Check for duplicates
print(f"NO. OF DUPLICATES IN DF: {df.duplicated().sum()}")
```

```
NO. OF MISSING VALLUES IN DF: 0
NO. OF DUPLICATES IN DF: 0
```

```
In [41]: # Check for unique values in each column
print("NO. OF UNIQUE VAL PER COL:")
for col in columns:
    print(f"\t{col}: {df[col].nunique()}")
```

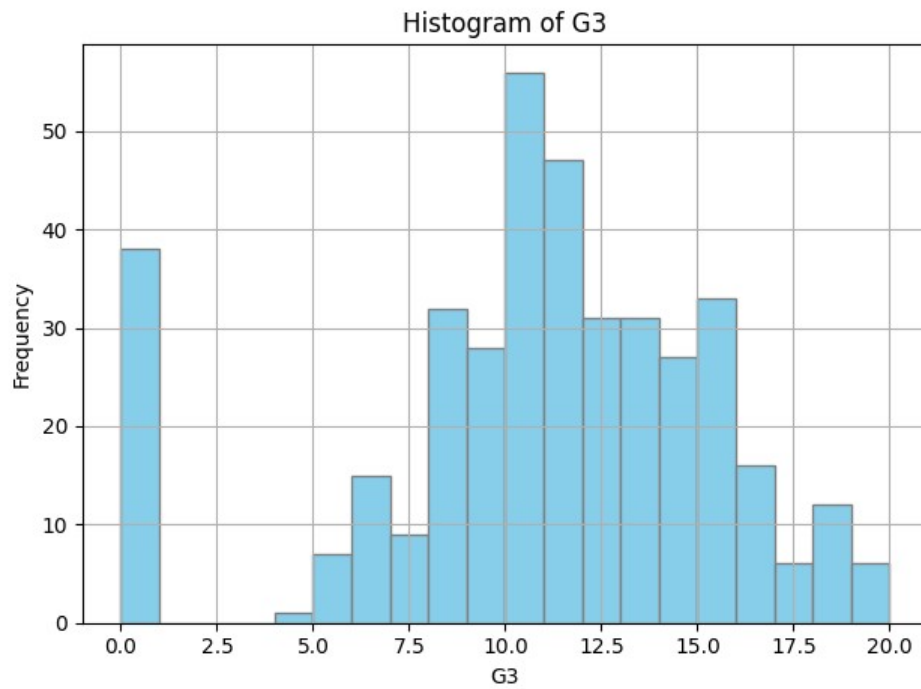
```
NO. OF UNIQUE VAL PER COL:
    school: 2
     sex: 2
    age: 8
address: 2
famsize: 2
Pstatus: 2
  Medu: 5
  Fedu: 5
  Mjob: 5
  Fjob: 5
reason: 4
guardian: 3
traveltime: 4
studytime: 4
failures: 4
schoolsup: 2
famsup: 2
paid: 2
activities: 2
nursery: 2
higher: 2
internet: 2
romantic: 2
famrel: 5
freetime: 5
goout: 5
Dalc: 5
Walc: 5
health: 5
absences: 34
    G1: 17
    G2: 17
    G3: 18
```

```
In [42]: # Show the numerical statistics of the data
df.describe()
```

```
Out[42]:
```

	age	Medu	Fedu	traveltime	studytime	failures	famrel	freetime	goout	Dalc	Walc	health	absences	G1
<b>count</b>	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00	395.00
<b>mean</b>	16.70	2.75	2.52	1.45	2.04	0.33	3.94	3.24	3.11	1.48	2.29	3.55	5.71	10.91
<b>std</b>	1.28	1.09	1.09	0.70	0.84	0.74	0.90	1.00	1.11	0.89	1.29	1.39	8.00	3.32
<b>min</b>	15.00	0.00	0.00	1.00	1.00	0.00	1.00	1.00	1.00	1.00	1.00	1.00	0.00	3.00
<b>25%</b>	16.00	2.00	2.00	1.00	1.00	0.00	4.00	3.00	2.00	1.00	1.00	3.00	0.00	8.00
<b>50%</b>	17.00	3.00	2.00	1.00	2.00	0.00	4.00	3.00	3.00	1.00	2.00	4.00	4.00	11.00
<b>75%</b>	18.00	4.00	3.00	2.00	2.00	0.00	5.00	4.00	4.00	2.00	3.00	5.00	8.00	13.00
<b>max</b>	22.00	4.00	4.00	4.00	4.00	3.00	5.00	5.00	5.00	5.00	5.00	5.00	75.00	19.00

```
In [43]: df['G3'].hist(bins=20, edgecolor='gray', color='skyblue')
plt.title('Histogram of G3')
plt.xlabel('G3')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```



```
In [44]: categorical_columns = df.columns[df.dtypes == 'object']
print(f"CATEGORICAL COLUMNS: \n{categorical_columns}")

CATEGORICAL COLUMNS:
Index(['school', 'sex', 'address', 'famsize', 'Pstatus', 'Mjob', 'Fjob',
       'reason', 'guardian', 'schoolsup', 'famsup', 'paid', 'activities',
       'nursery', 'higher', 'internet', 'romantic'],
      dtype='object')
```

```

In [45]: n = len(categorical_columns)
ncols = 4
nrows = n // ncols if n % ncols == 0 else n // ncols + 1

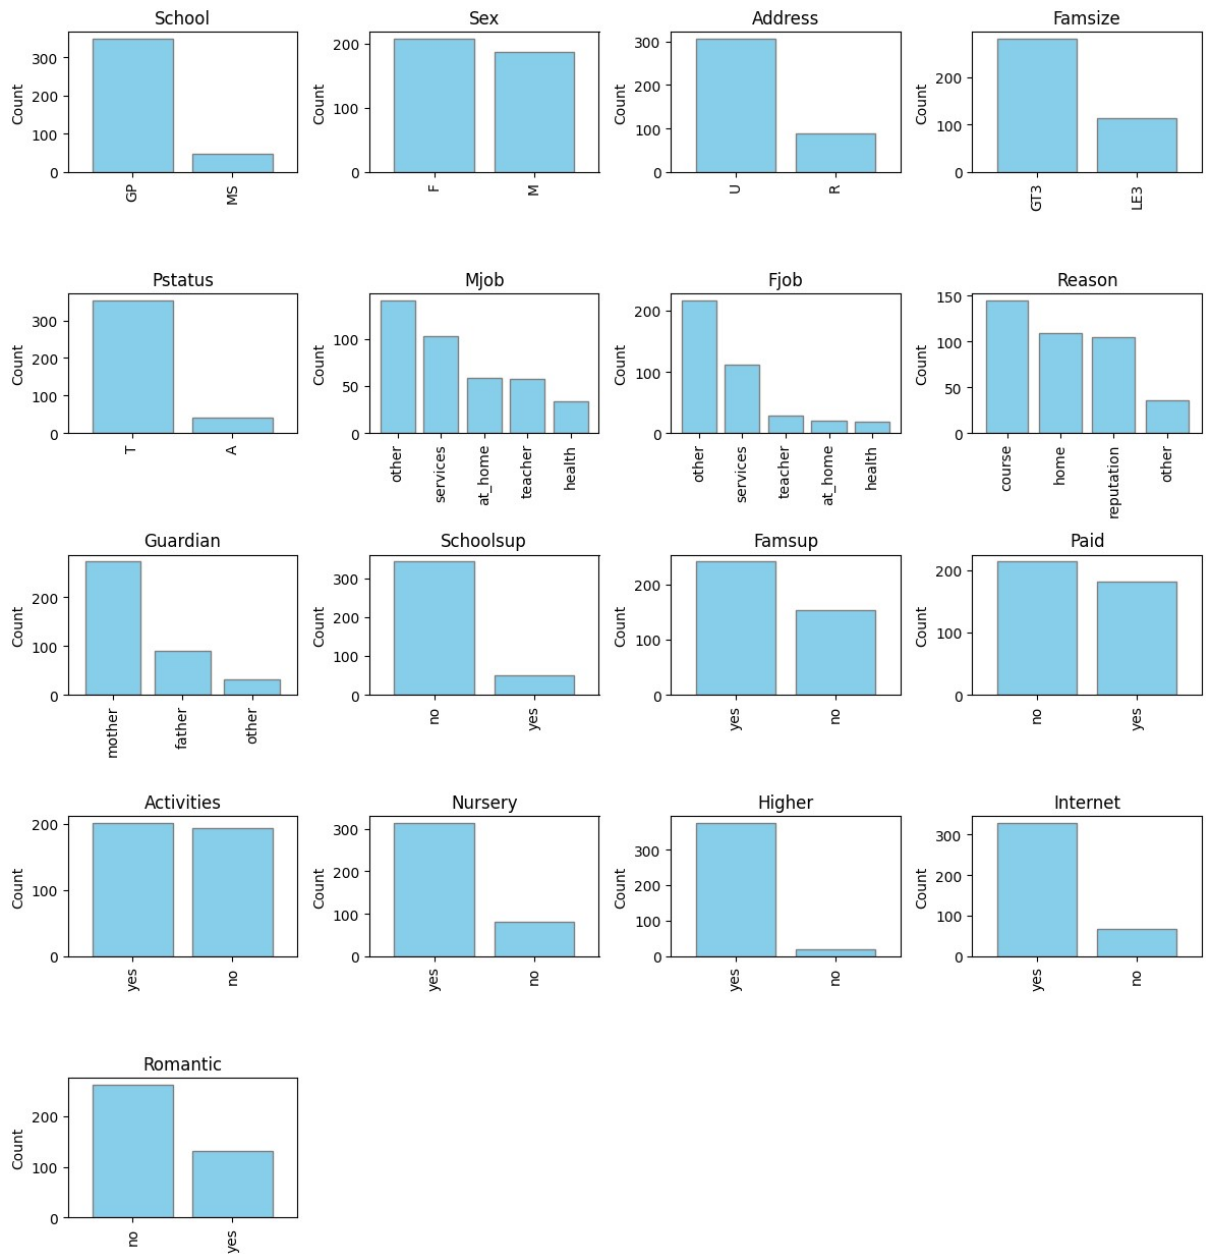
fig, axs = plt.subplots(nrows, ncols, figsize=(3*ncols, 2.5*nrows))

axs = axs.flatten()
for ax in axs[n:]:
    fig.delaxes(ax)

for ax, column in zip(axs, categorical_columns):
    value_counts = df[column].value_counts(ascending=False)
    value_counts.plot(kind='bar', ax=ax, width=0.8, color='skyblue', edgecolor='gray')
    ax.set_title(f'{column.capitalize()}')
    ax.set_ylabel('Count')

plt.tight_layout()
plt.show()

```



```
In [46]: num_plots = len(categorical_columns)
num_cols = 4
num_rows = num_plots // num_cols if num_plots % num_cols == 0 else num_plots // num_cols + 1

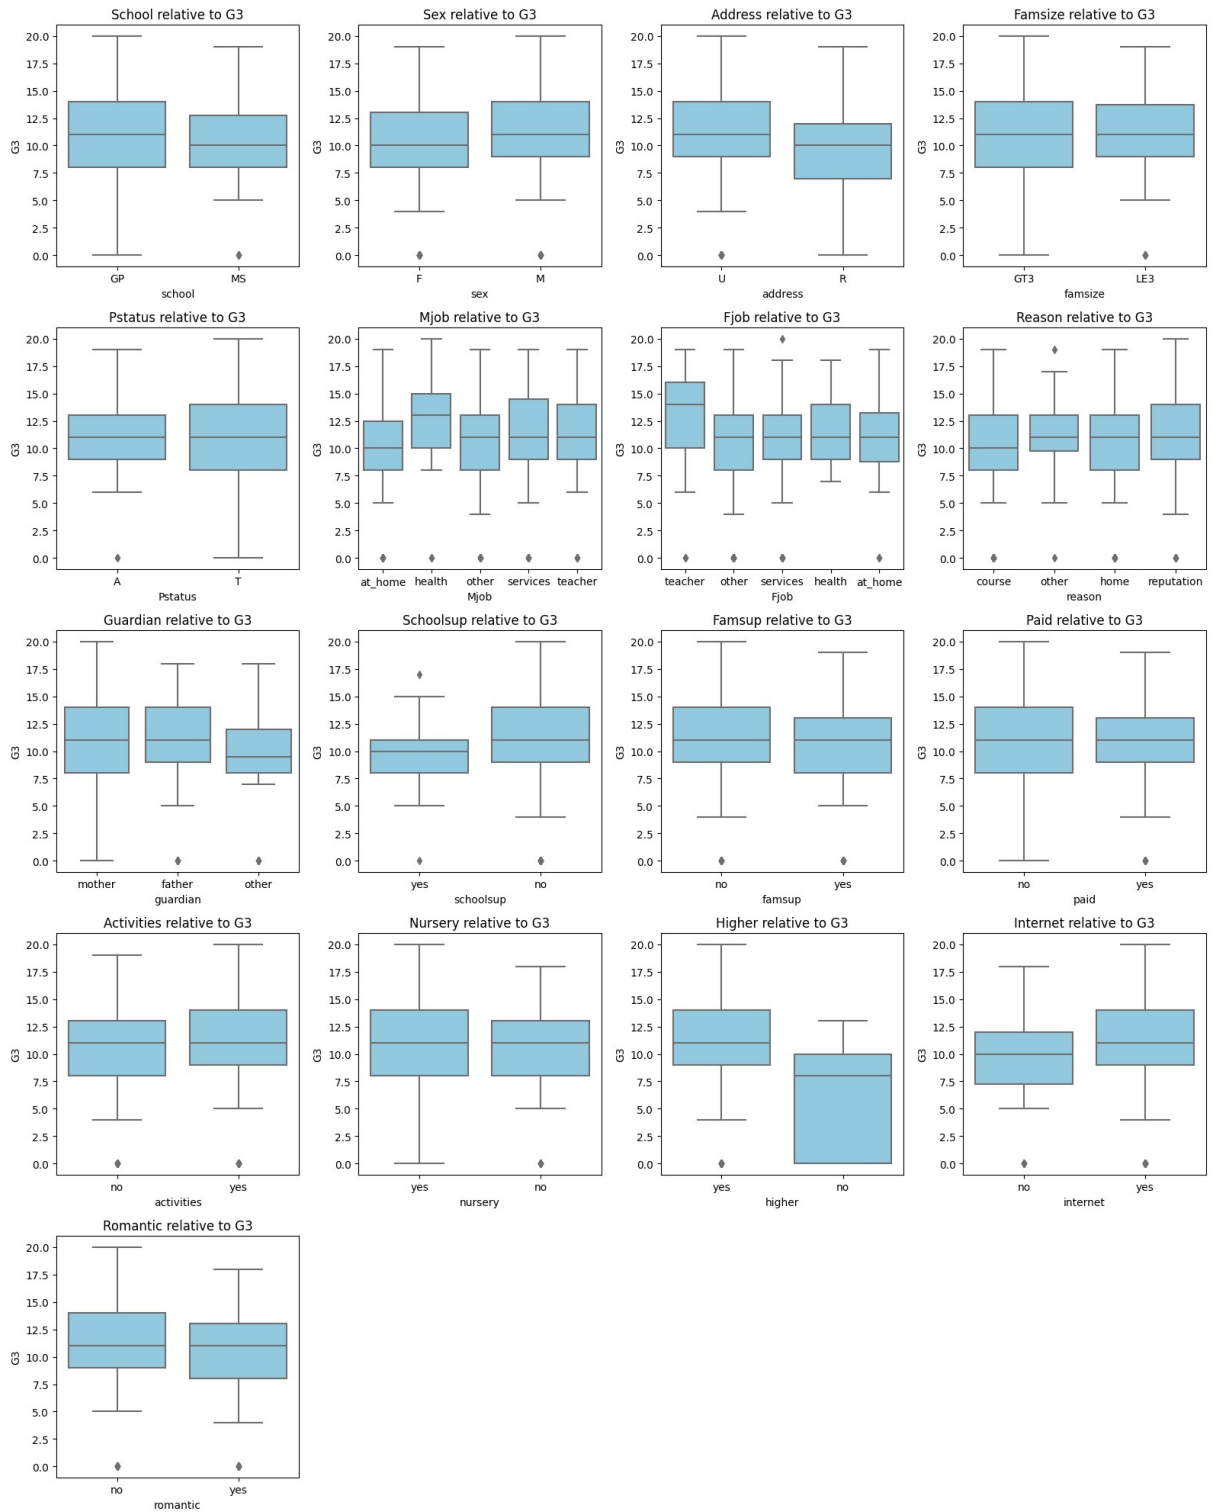
fig, axs = plt.subplots(num_rows, num_cols, figsize=(num_cols*4, num_rows*4))

for i, column in enumerate(categorical_columns):
    row = i // num_cols
    col = i % num_cols
    sns.boxplot(x=column, y='G3', data=df, ax=axs[row, col], color='skyblue')
    axs[row, col].set_title(f'{column.capitalize()} relative to G3')

# Remove empty subplots
if num_plots % num_cols != 0:
    for col in range(num_plots % num_cols, num_cols):
        fig.delaxes(axs[num_rows-1, col])

plt.tight_layout()
plt.show()
```





```
In [47]: numerical_columns = df.columns[df.dtypes != 'object']
numerical_columns = numerical_columns.drop('G3')
print(f"NUMERICAL COLUMNS: \n{numerical_columns}")
```

```
NUMERICAL COLUMNS:
Index(['age', 'Medu', 'Fedu', 'traveltime', 'studytime', 'failures', 'famrel',
      'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences', 'G1', 'G2'],
      dtype='object')
```

```

In [48]: n = len(numerical_columns)
ncols = 4
nrows = n // ncols if n % ncols == 0 else n // ncols + 1

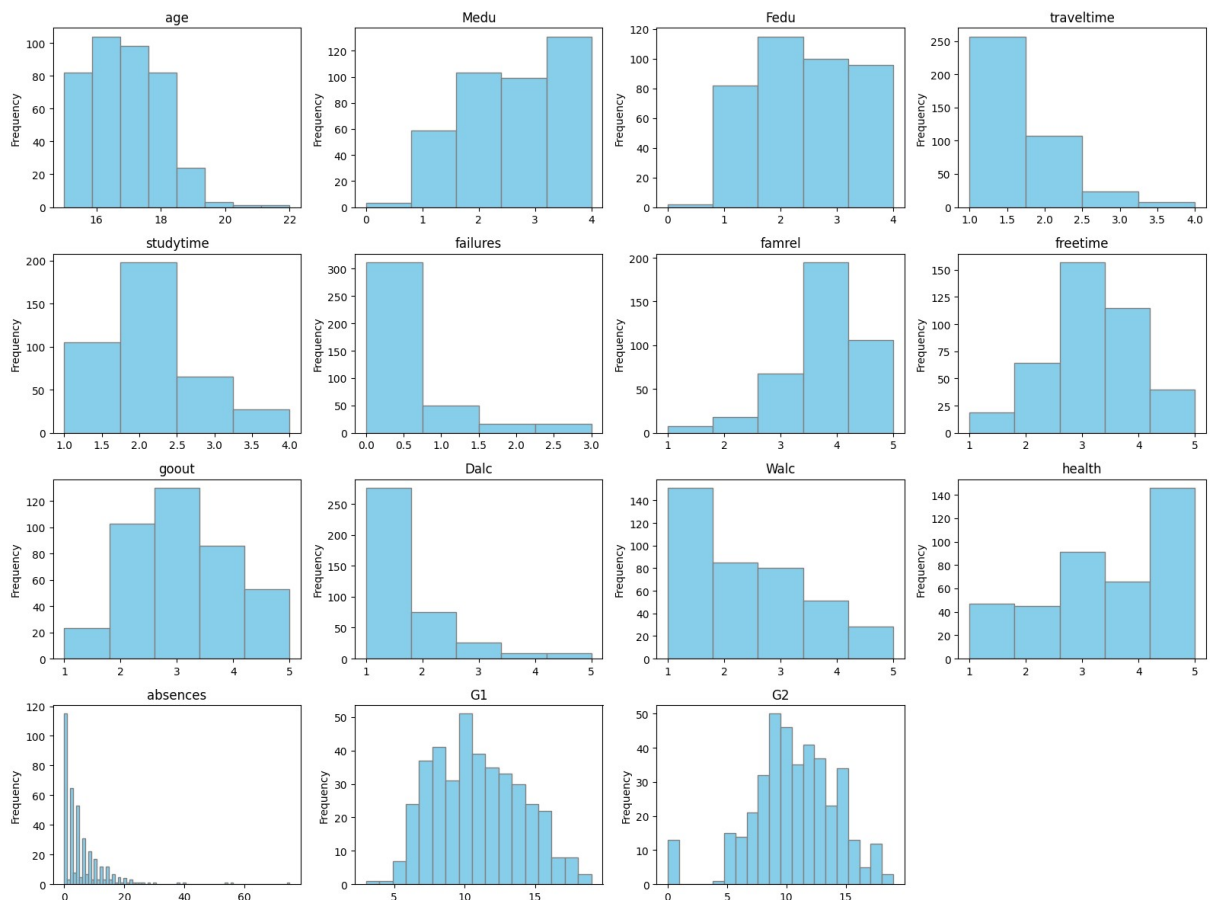
fig, axs = plt.subplots(nrows, ncols, figsize=(4*ncols, 3*nrows))

axs = axs.flatten()
for ax in axs[n:]:
    fig.delaxes(ax)

for ax, column in zip(axs, numerical_columns):
    bins = df[column].max() - df[column].min() + 1
    ax.hist(df[column], bins=int(bins), color='skyblue', edgecolor='gray')
    ax.set_title(f'{column}')
    ax.set_xlabel('')
    ax.set_ylabel('Frequency')

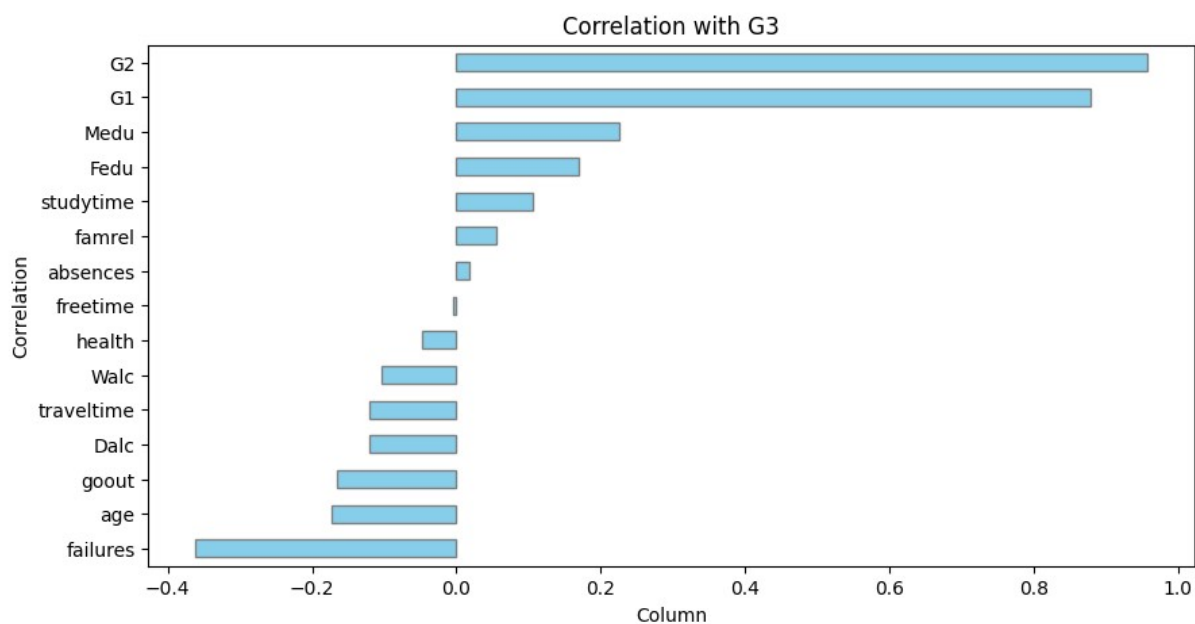
plt.tight_layout()
plt.show()

```



```
In [49]: correlations = df[numerical_columns].corrwith(df['G3'], method='spearman')

plt.figure(figsize=(10, 5))
correlations.sort_values().plot(kind='barh', color='skyblue', edgecolor='gray')
plt.title('Correlation with G3')
plt.xlabel('Column')
plt.ylabel('Correlation')
plt.show()
```



## ▼ 1.0.2 Model Development

```
In [50]: X = df.drop('G3', axis=1)
y = df['G3']

print(f"Feature Shape: {X.shape}")
print(f"Target Shape: {y.shape}")
```

Feature Shape: (395, 32)  
Target Shape: (395,)

```
In [51]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123, shuffle=True)

print(f"X Train Shape: {X_train.shape}")
print(f"Y Train Shape: {y_train.shape}")
print(f"X Test Shape: {X_test.shape}")
print(f"Y Test Shape: {y_test.shape}")
```

X Train Shape: (316, 32)  
Y Train Shape: (316,)  
X Test Shape: (79, 32)  
Y Test Shape: (79,)

```
In [52]: # Encode categorical variables
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OrdinalEncoder(), categorical_columns),
        ('num', StandardScaler(), numerical_columns)
    ])

X_train = preprocessor.fit_transform(X_train)
X_test = preprocessor.transform(X_test)
```

```
In [53]: reg = LazyRegressor(ignore_warnings=False, custom_metric=None)
models, predictions = reg.fit(X_train, X_test, y_train, y_test)
```

```
29%|███████| 12/42 [00:00<00:01, 18.05it/s]
```

GammaRegressor model failed to execute

Some value(s) of y are out of the valid range of the loss 'HalfGammaLoss'.

```
100%|██████████| 42/42 [00:03<00:00, 13.03it/s]
```

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000050 seconds.

You can set `force\_row\_wise=true` to remove the overhead.

And if memory is not enough, you can set `force\_col\_wise=true`.

[LightGBM] [Info] Total Bins 182

[LightGBM] [Info] Number of data points in the train set: 316, number of used features: 31

[LightGBM] [Info] Start training from score 10.262658

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

```
In [54]: sorted_model = models.sort_values('R-Squared', ascending=False)
best_model = sorted_model.iloc[0]
print(f"BEST MODEL: {best_model.name}")
print(best_model)
print("\nSORTED MODELS BY R-SQUARED:")
print(sorted_model)
```

```
BEST MODEL: RandomForestRegressor
Adjusted R-Squared    0.87
R-Squared            0.92
RMSE                 1.34
Time Taken           0.33
Name: RandomForestRegressor, dtype: float64
```

SORTED MODELS BY R-SQUARED:

	Adjusted R-Squared	R-Squared	RMSE	Time Taken
Model				
RandomForestRegressor	0.87	0.92	1.34	0.33
LGBMRegressor	0.86	0.92	1.35	0.03
HistGradientBoostingRegressor	0.86	0.92	1.37	0.22
GradientBoostingRegressor	0.86	0.91	1.39	0.11
DecisionTreeRegressor	0.85	0.91	1.41	0.01
BaggingRegressor	0.85	0.91	1.44	0.05
XGBRegressor	0.84	0.91	1.45	0.07
ExtraTreesRegressor	0.80	0.88	1.64	0.31
AdaBoostRegressor	0.79	0.88	1.66	0.04
OrthogonalMatchingPursuit	0.74	0.85	1.85	0.02
LassoLarsIC	0.73	0.84	1.88	0.02
OrthogonalMatchingPursuitCV	0.73	0.84	1.88	0.00
LarsCV	0.73	0.84	1.89	0.03
LassoLarsCV	0.73	0.84	1.89	0.02
LassoCV	0.73	0.84	1.90	0.08
ElasticNetCV	0.73	0.84	1.91	0.07
RidgeCV	0.72	0.84	1.93	0.00
BayesianRidge	0.72	0.84	1.93	0.02
Ridge	0.72	0.83	1.94	0.02
TransformedTargetRegressor	0.72	0.83	1.94	0.00
LinearRegression	0.72	0.83	1.94	0.02
Lars	0.72	0.83	1.94	0.00
SGDRegressor	0.72	0.83	1.94	0.01
HuberRegressor	0.72	0.83	1.95	0.02
RANSACRegressor	0.71	0.83	1.96	0.11
LinearSVR	0.70	0.83	1.99	0.01
ExtraTreeRegressor	0.66	0.80	2.14	0.01
PoissonRegressor	0.63	0.78	2.23	0.02
Lasso	0.60	0.76	2.31	0.02
LassoLars	0.60	0.76	2.31	0.01
TweedieRegressor	0.56	0.74	2.43	0.02
MLPRegressor	0.55	0.73	2.46	0.25
ElasticNet	0.54	0.73	2.48	0.00
SVR	0.45	0.68	2.71	0.02
NuSVR	0.43	0.66	2.76	0.00
KNeighborsRegressor	0.27	0.57	3.13	0.02
PassiveAggressiveRegressor	0.12	0.48	3.44	0.02
QuantileRegressor	-0.70	-0.00	4.76	1.16
DummyRegressor	-0.74	-0.03	4.82	0.01
KernelRidge	-7.34	-3.92	10.55	0.02
GaussianProcessRegressor	-9.66	-5.29	11.94	0.03

```
In [55]: best_model_ = reg.models[best_model.name]
model = best_model_.named_steps['regressor']
best_model_
```

```
Out[55]: Pipeline(steps=[('preprocessor',
                          ColumnTransformer(transformers=[('numeric',
                                                            Pipeline(steps=[('imputer',
                                                                 SimpleImputer()),
                                                                 ('scaler',
                                                                  StandardScaler()))],
                                                            Int64Index([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10
, 11, 12, 13, 14, 15, 16,
                          17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31],
                          dtype='int64')),
                          ('categorical_low',
                           Pipeline(steps=[('imputer',
                                                                 SimpleImputer(fill_value='missing',
                                                                 strategy='constant')),
                           ('encoding',
                            OneHotEncoder(handle_unknown='ignore',
                                             sparse=False))]),
                          Int64Index([], dtype='int64')),
                          ('categorical_high',
                           Pipeline(steps=[('imputer',
                                                                 SimpleImputer(fill_value='missing',
                                                                 strategy='constant')),
                           ('encoding',
                            OrdinalEncoder())]),
                          Int64Index([], dtype='int64'))]),
                          ('regressor', RandomForestRegressor(random_state=42))])
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [56]: X_column = pd.DataFrame(X_test, columns=X.columns)
explainer = shap.Explainer(model.predict, X_column)

shap_values = explainer.shap_values(X_column)
num_features = X_column.shape[1]
```

```
PermutationExplainer explainer: 80it [00:11, 1.05it/s]
```

```
In [57]: shap.summary_plot(shap_values, X_column, plot_type="bar", max_display=num_features)
```





school

0.0

0.5

1.0

1.5

2.0

2.5

3.0

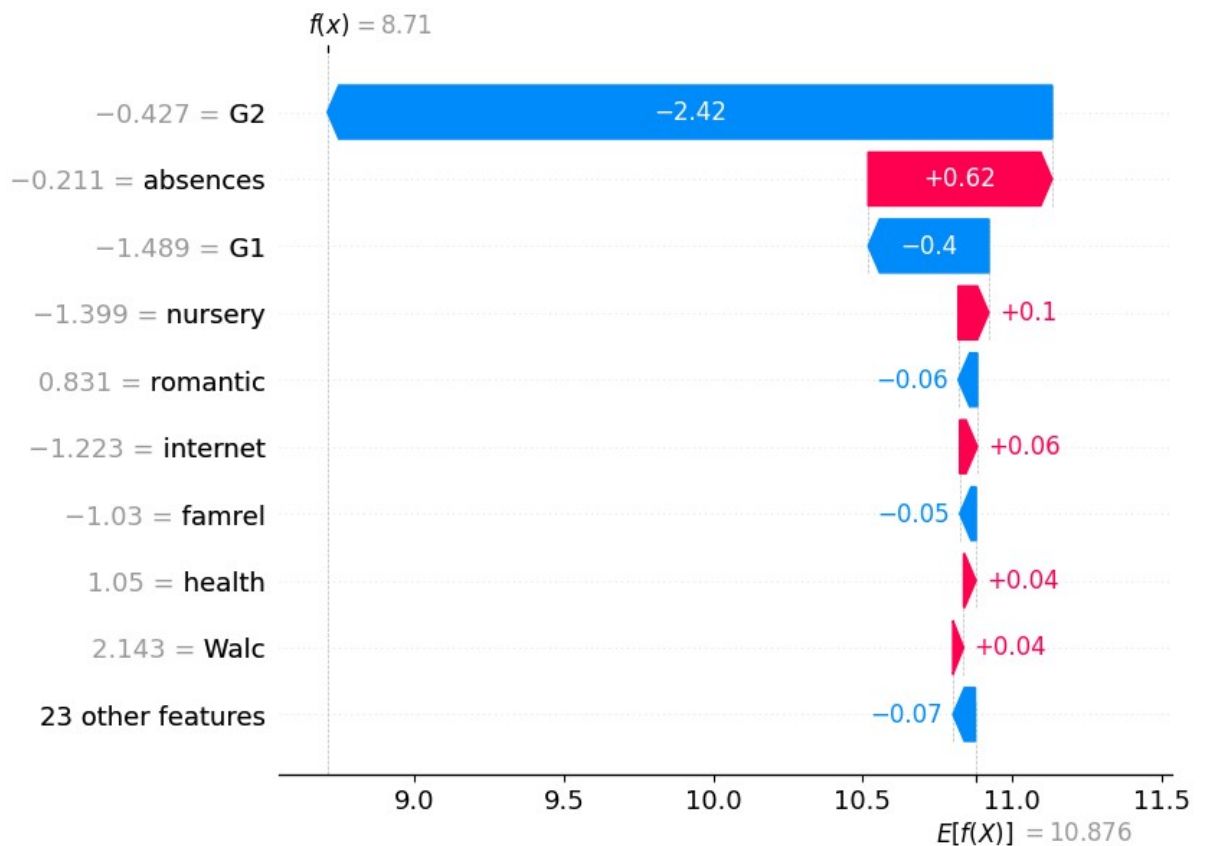
mean(|SHAP value|) (average impact on model output magnitude)

```
In [65]: sv = explainer(X_column)
y_test_pred = model.predict(X_test)
y_test_df = y_test.to_frame() # Convert y_test to DataFrame


for student_no in [0,1,3]:
    print(f"\nSTUDENT LOC: {student_no+1}")
    print(f"Predicted Performance: {y_test_pred[student_no]}")
    print(f"Actual Performance: {y_test_df['G3'].iloc[student_no]}") # Use iloc for DataFrame
    exp = shap.Explanation(sv.values[student_no,:],
                        sv.base_values[student_no],
                        data=X_column.values[student_no,:],
                        feature_names=X_column.columns)

    shap.plots.waterfall(exp)
```

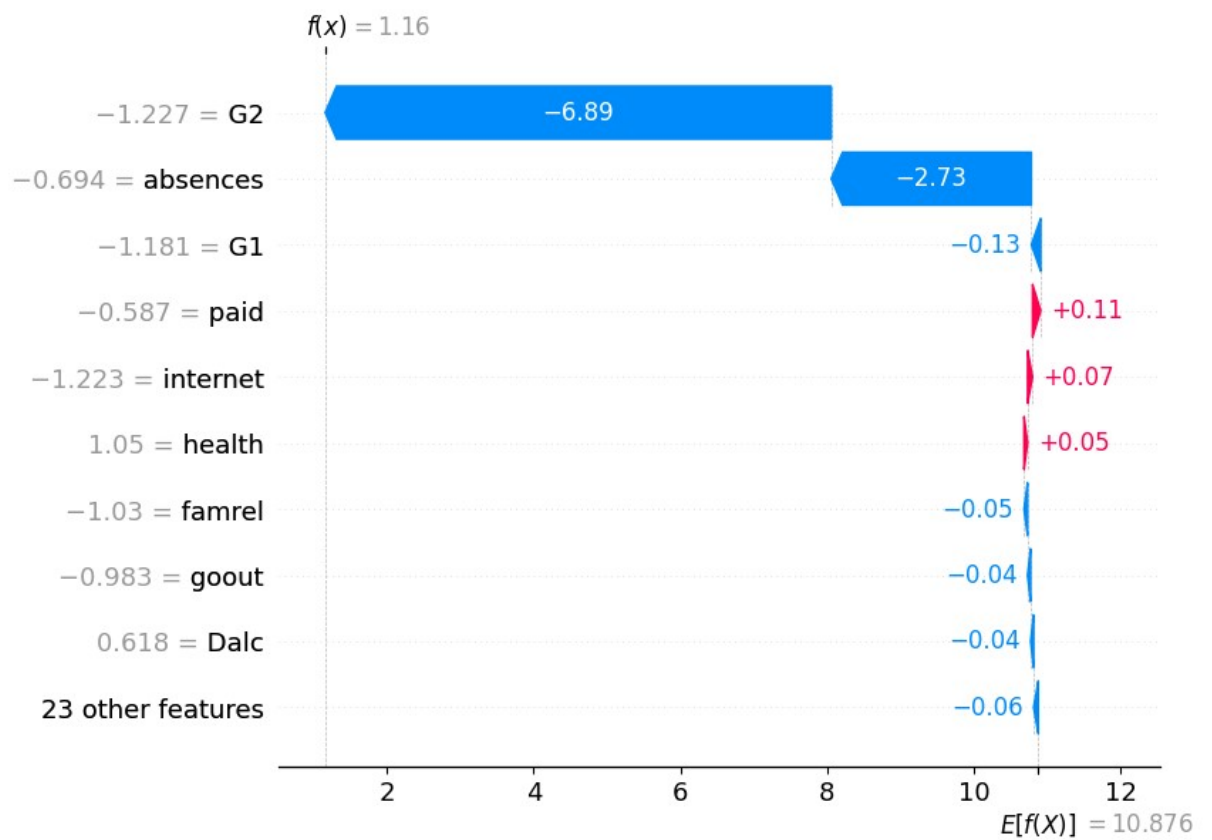
STUDENT LOC: 1  
Predicted Performance: 8.71  
Actual Performance: 8



STUDENT LOC: 2  
Predicted Performance: 13.15  
Actual Performance: 13



STUDENT LOC: 4  
 Predicted Performance: 1.16  
 Actual Performance: 0



```
In [66]: print(f"BEST MODEL: {best_model.name}")
print(best_model)
```

```
BEST MODEL: RandomForestRegressor
Adjusted R-Squared    0.87
R-Squared            0.92
RMSE                 1.34
Time Taken           0.33
Name: RandomForestRegressor, dtype: float64
```

### ▼ 1.0.3 INSIGHTS

#### Best Model based on LazyRegressor

Random Forest Regressor

Adjusted R-Squared = 0.87

R-Squared = 0.92

RMSE = 1.34

Time Taken = 0.33

Lower RMSE than the best performing model in the article for Mathematics data.

#### What are SHAP (Shapley Additive exPlanations) values?

SHAP values are a game-theoretic approach to explain the output of any machine learning model by fairly allocating credit for the model's output prediction ( $f_x$ ) among input features. Note that the magnitude of a SHAP value measures the contribution of a corresponding feature to the model's prediction, and the sign of a SHAP value signals whether the corresponding feature made a positive contribution (positive sign) or a negative contribution (negative sign). Adding all SHAP values to the base value or the average model prediction results in obtaining the model's output for an instance.

#### Feature Importance Plot (Bar Chart)

The figure shows the mean absolute SHAP value, which is the average impact of each feature on the predictions for the final grade (G3) over all students. The most significant feature in predicting the final grade is the second-period grade (G2). This is followed by the number of absences, suggesting that higher absences negatively affect the final grade. The first-period grade (G1) is also essential but relatively less necessary than G2.

#### Individual SHAP Value Plots

The figure is a collection of SHAP value plots for the feature values of a particular instance. A SHAP value plot shows the contribution of the feature values on the prediction of the final grade (G3) of a particular student. The base value is 10.876, and each plot is for a different student. Here are some observations for 3 students:

Student 1:

- G2 strongly negatively influences the predicted grade by -2.41.
- Number of absences positively influence the grade by 0.63, meaning that the student may have a low number of absences.
- G1 negatively influences the grade, but less than G2.
- Predicted G3 is 8.71. True G3 is 8.

Student 2:

- G2 strongly positively influences the predicted grade by 2.03.
- Number of absences has a positive impact of 0.27.
- Weekend Alcohol Consumption (Walc) has a slight negative impact of -0.15.
- Predicted G3 is 13.15. True G3 is 13.

Student 4:

- G2 has a significant negative impact of -6.88.
- Number of absences strongly negatively influences the predicted grade by -2.62.
- G1 has a slight negative impact of -0.16.
- Predicted G3 is 1.16. True G3 is 0.

In summary, based on the SHAP values, the most important factors are grades from prior periods (G2, G1) and the number of absences, while other features have varying levels of influence depending on the particular student's data.